

OS PROJECT 2

PROCESS SCHEDULING IN LINUX

Advisor: Chien-Chung Ho

TA: Zong-Zhe Yang

Shih-Hao Chin



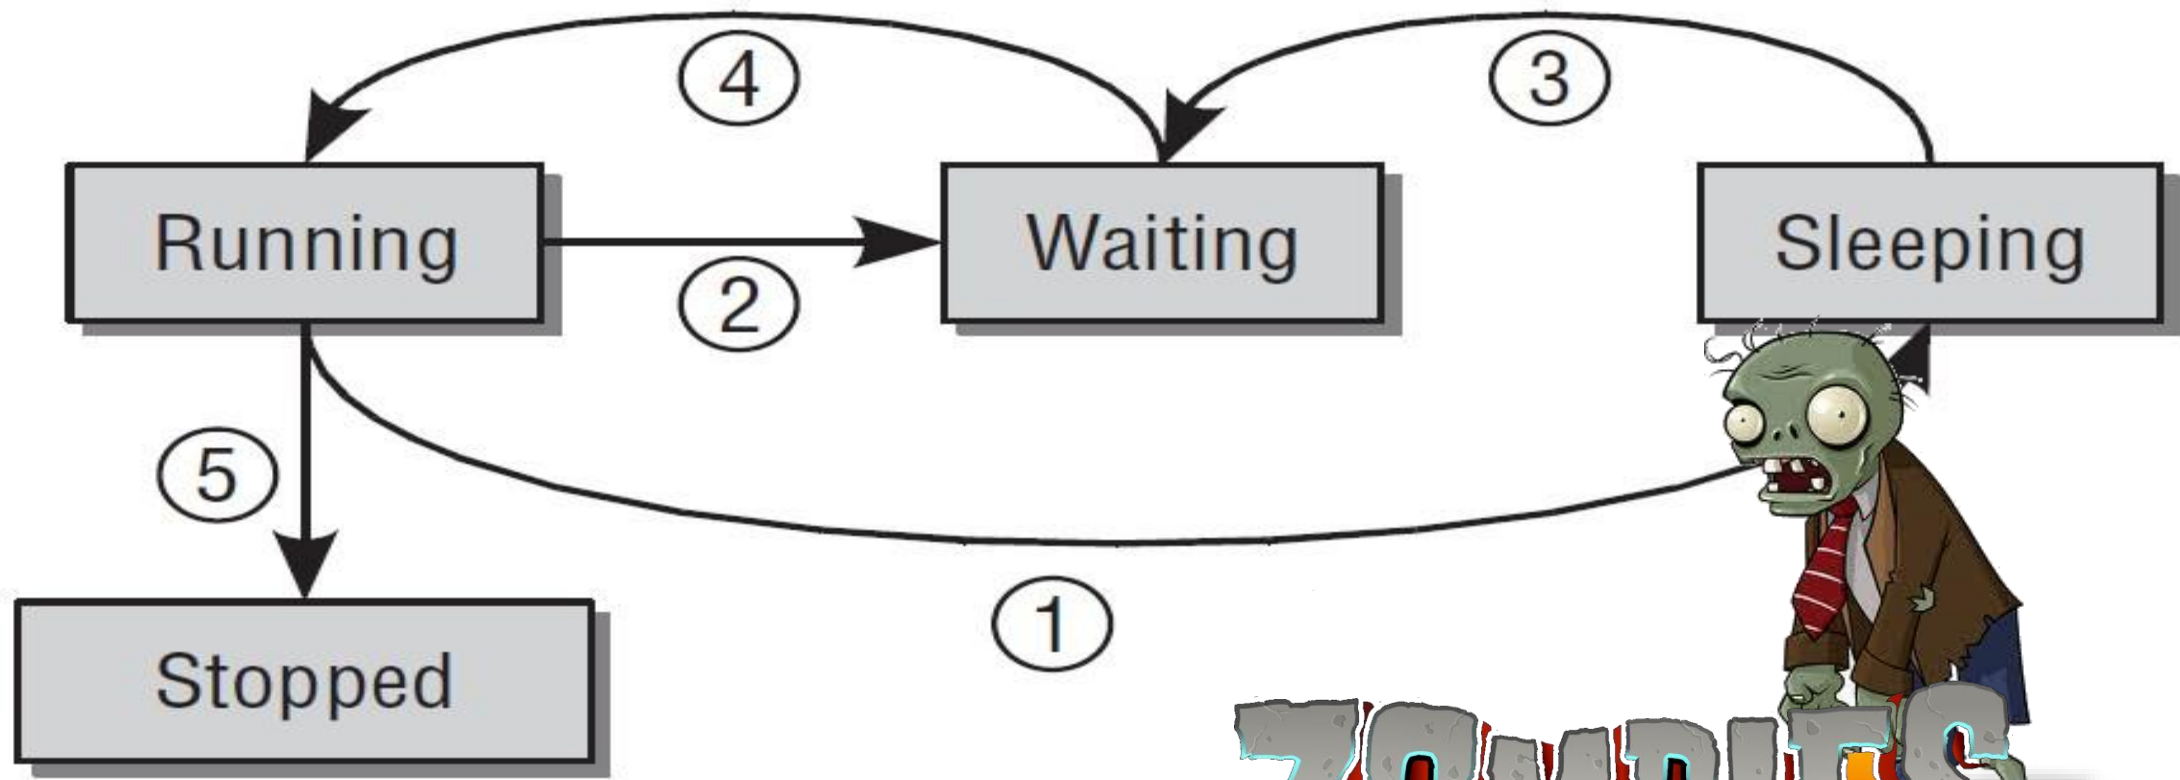
OUTLINE

- Introduction
- Project Requirements
- Submission Rules
- References

PROCESS LIFE CYCLE

- A process is *not* always ready to run.
- The scheduler must know the status of every
 - process in the system when switching between tasks.
- A process may have one of the following states:
 - **Running** — The process is executing at the moment.
 - **Waiting** — The process is able to run but is not allowed to because the CPU is allocated to another process. The scheduler can select the process at the next task switch.
 - **Sleeping** — The process is sleeping and cannot run because it is waiting for an external event. The scheduler *cannot* select the process at the next task switch.
- The system saves all processes in a process table.

TRANSITIONS BETWEEN PROCESS STATES



SCHEDULING IN LINUX (1/2)

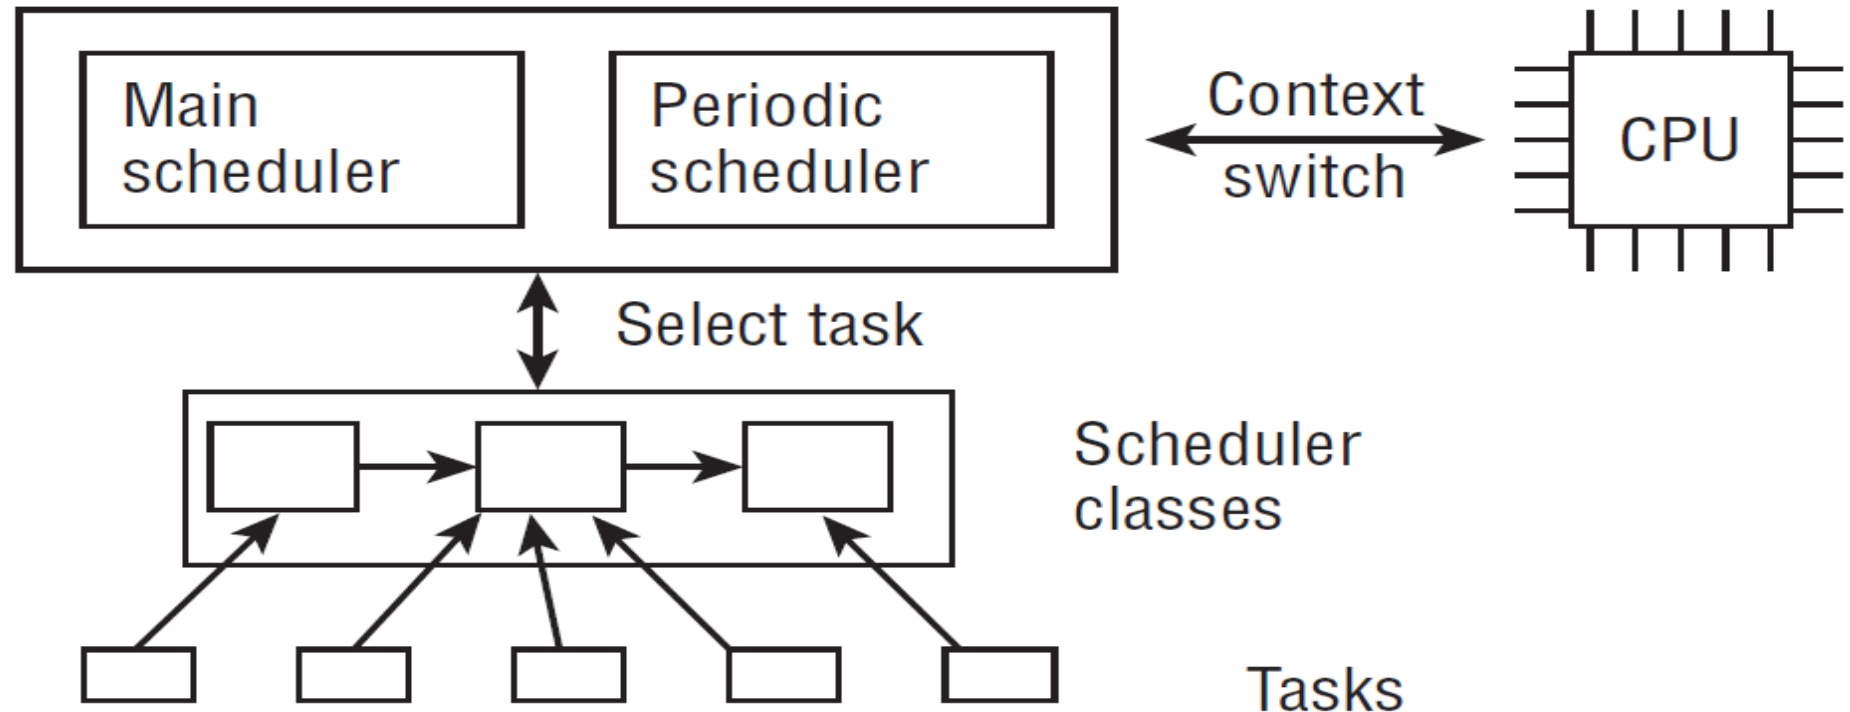

- The schedule function is the starting point to an understanding of scheduling operations.
- It is defined in “`kernel/sched/core.c`” and is one of the most frequently invoked functions in the kernel code.
- Not only *priority scheduling* but also two other soft real-time policies required by the POSIX standard are implemented.
 - E.g., completely fair scheduling, real-time scheduling and scheduling of the idle task, *etc.*

SCHEDULING IN LINUX (2/2)

- The scheduler uses a series of data structures to sort and manage the processes in the system.
- Scheduling can be activated in two ways:
 - **Main scheduler:** Either directly if a task goes to sleep or wants to yield the CPU for other reasons,
 - **Periodic scheduler:** Or by a periodic mechanism that is run with constant frequency to check from time to time if switching tasks is necessary

Generic scheduler = Main + Periodic schedulers

OVERVIEW OF THE COMPONENTS OF THE SCHEDULING SUBSYSTEM



EXAMPLE:

PROCESS REPRESENTATION IN LINUX

- In Linux, all concerned with processes and programs are built around a data structure: **task_struct**.

```
struct task_struct {  
    volatile long state;      /* -1 unrunnable, 0 runnable, >0 stopped */  
    void *stack;  
    atomic_t usage;  
    unsigned int flags;       /* per process flags, defined below */  
    unsigned int ptrace;  
  
#ifdef CONFIG_SMP  
    struct llist_node wake_entry;  
    int on_cpu;  
    unsigned int wakee_flips;  
    unsigned long wakee_flip_decay_ts;  
    struct task_struct *last_wakee;  
  
    int wake_cpu;  
#endif
```

... see more in "include/linux/sched.h"

HOW TO DESIGNATE A SCHEDULER FOR TASKS?

```
int prio, static_prio, normal_prio;
unsigned int rt_priority;
const struct sched_class *sched_class;
struct sched_entity se;
struct sched_rt_entity rt;
#ifdef CONFIG_CGROUP_SCHED
struct task_group *sched_task_group;
#endif
struct sched_dl_entity dl;

#ifdef CONFIG_PREEMPT_NOTIFIERS
/* list of struct preempt_notifier: */
struct hlist_head preempt_notifiers;
#endif

#ifdef CONFIG_BLK_DEV_IO_TRACE
unsigned int btrace_seq;
#endif

unsigned int policy;
int nr_cpus_allowed;
cpumask_t cpus_allowed;
```



OUTLINE

- Introduction
- **Project Requirements**
- Submission Rules
- References

LINUX SCHEDULING POLICIES

- Linux Scheduling Policies
 - Normal Scheduling policies (Non-real-time)
 - SCHED_OTHER, SCHED_BATCH, SCHED_IDLE.
 - Real-Time policies
 - SCHED_FIFO, SCHED_RR.
- The default scheduling policy is **non-real-time**.
- In this part, using Linux **real-time scheduling policy (FIFO)** to schedule threads in a process.

PROJECT REQUIREMENTS

- 1) Write a C program (`sched_test.c`) to create threads
- 2) Need busy-waiting to ensure that process exist long
- 3) Run the program by default time sharing schedule policy and show the result. (20%)
 - Ex: `sudo ./sched_test 5`
- 4) Run the program by real time scheduling policy (RR,FIFO) and show the result.(40%)
 - Ex: `sudo ./sched_test 5 SCHED_FIFO`
- 5) The parameters must be able to select the number of threads and the policy

HINT

- Set CPU affinity
- sched_setscheduler
- Set the priority of real time process sched_param *param
- The permission to run real time process

PROJECT REQUIREMENTS

- 6) Install **kernelShark** + **trace-cmd** to analyze process (ex:FIFO)

<https://hackmd.io/@ULxjDFy0QLKSAQnD5DMKtQ/HJfBkEYnH>

[illegible]

PROJECT REQUIREMENTS

- Report(.PDF) (60%)
 - What do you learn from this project
 - Is there any problem, state it Your implementation details
 - The test results (contains the resulting diagram generated by kernel shark)
 - At most 6 pages



OUTLINE

- Introduction
- Project Requirements
- Submission Rules
- References


SUBMISSION RULES

- Be packed as one file named “OSPJ2_Team##.tar”
 - OSPJ2_Team## (directory)
 - report.pdf
 - sched_test.c
 - sched_test.o

SUBMISSION RULES

- Project deadline: 2019/12/10 (23:59)
 - Delayed submissions yield 5 point deduction per day
- Upload your team project to the FTP site.
 - FTP server: 140.123.105.185
 - user ID: OS2019 ; pass: 2019OS
- If you've submitted files for many times, TAs will only check the up-to-date version.
- If the up-to-date file misses the deadline, it will be regarded as the delayed submission.
- **DO NOT COPY THE HOMEWORK**

CONTACT TAs

- If you have any problem about the projects or this course, you can contact TAs by the following ways.
- Facebook: CCU OS2019 Fall Group 
 - <https://www.facebook.com/groups/1319131934920970/>
 - E-mail: Shih-Hao Chin: ae8681239@gmail.com
Zong-Zhe Yang : daankfust123@gmail.com



OUTLINE

- Introduction
- Project Requirements
- Submission Rules
- References

REFERENCES



- Reference Book
 - Professional Linux® Kernel Architecture, Wolfgang Mauerer, Wiley Publishing, Inc.
- Process Scheduling
 - <http://www.cs.rutgers.edu/~pxk/416/notes/07-scheduling.html>